

Stathis Galanakis

GENERATIVE AI | 3D VISION | COMPUTER VISION SYSTEMS

London, UK

@stathisgalan@gmail.com

in stathisgln |  Stathis Galanakis |  stathisgln.github.io |  stathisgln

Summary

Computer Vision Researcher (PhD, Imperial College London) specialising in generative modelling and 3D human reconstruction. My work focuses on diffusion-based approaches for monocular 3D reconstruction and multi-view consistency, with first-author publications at ICCV and WACV. I am interested in bridging generative models with structured 3D representations under low-data and real-world constraints.

Experience

Research Associate in Generative Computer Vision Models

London, UK

Department of Computing, Imperial College London

March 2025 - Present

- Investigated diffusion-based synthetic data generation for medical imaging under low-data regimes, resulting in an 8% improvement in pathology classification through image-text dataset curation (500k pairs).
- Explored the use of large language models (LLMs) for synthetic medical caption generation and automated classification of pathology reports.
- Supervise Master's students on research projects in generative AI and world modelling.

Computer Vision Internship

London, UK

Huawei UK

January 2022 - March 2025

- Proposed diffusion and NeRF-based approaches for monocular 3D face reconstruction under geometric ambiguity, achieving state-of-the-art identity preservation performance.
- Developed synthetic identity datasets (10k+ identities) to improve robustness to pose and appearance variation.
- Designed a 3D Gaussian Splatting framework for animatable head avatars, enabling consistent identity generation across viewpoints.

Research Assistant

London, UK

Business School, Imperial College London

February 2021 - January 2022

- Applied machine learning to multimodal satellite imagery and weather data for crop yield prediction under spatial data sparsity.
- Developed synthetic data augmentation strategies to improve generalisation in underrepresented geographic regions.

Computer Vision Scientist

London, UK

ArielAI (acquired by Snap Inc.)

January 2020 - September 2020

- Developed large-scale data collection and filtering pipelines for training vision models, incorporating human-in-the-loop annotation to improve dataset quality.
- Contributed to scalable data infrastructure supporting model development in a startup environment later acquired by Snap Inc.

R&D, ML Engineer

Athens, Greece

Pobuca Ltd

May 2018 - January 2019

- Designed and deployed an end-to-end product detection system for retail shelf images (25+ classes), integrating semi-automated annotation and training workflows.

Selected Publications (7 total)

DERMAFLUX: Synthetic Skin Lesion Generation with Rectified Flows for Enhanced Image Classification
Galanakis S, Koliouris A, and Zafeiriou S. Under peer review
DERMAFLUX is a rectified flow-based text-to-image model for synthetic skin lesion image synthesis.

STARCaster: Spatio-Temporal AutoRegressive Video Diffusion for Identity- and View-Aware Talking Portraits
Paraperas F, Galanakis S, Potamias R, Kainz B, and Zafeiriou S. ICML 2026

SpinMeRound: Consistent Multi-View Identity Generation Using Diffusion Models
Galanakis S, Lattas A, Moschoglou S, Kainz B, and Zafeiriou S. ICCV 2025
SpinMeRound is a diffusion-based approach for a multi-view identity-consistent generation from a single facial image.

ImHead: A Large-scale Implicit Morphable Model for Localized Head Modeling
Potamias R, Galanakis S, Deng J, Papaioannou A, and Zafeiriou S. ICCV 2025

FitDiff: Robust monocular 3D facial shape and reflectance estimation using Diffusion Models
Galanakis S, Lattas A, Moschoglou S, and Zafeiriou S. WACV 2025
FitDiff is a multi-modal diffusion model, which jointly models facial geometry and appearance.

3DMM-RF: Convolutional Radiance Fields for 3D Face Modeling
Galanakis S, Gecer B, Lattas A, and Zafeiriou S. WACV 2023
3DMM-RF is an implicit 3D Morphable model that can accurately model a subject's identity, pose and expression under arbitrary illumination, by utilizing a style-based generator.

Education

PhD in Computer Science Imperial College London, UK
Supervisor: Prof. Stefanos Zafeiriou April 2025
Thesis: Advancements in Face Reconstruction from a Single Image.
Research focused on diffusion and generative modeling for monocular 3D face reconstruction, improving identity preservation and geometric stability under in-the-wild conditions. Work resulted in peer-reviewed publications (ICCV, WACV), including first-author contributions.

Diploma/M.Eng. in Electrical and Computer Engineering NTUA, Greece
Supervisor: Prof. Petros Maragos November 2019
Thesis: Human Action Recognition and Localisation in Videos.

Technical Strengths

Programming:	Python, CUDA, C++, Claude (LLM-assisted development)
Frameworks:	PyTorch, PyTorch Lightning, PyTorch3D, Accelerate
Research Areas:	GANs, Diffusion Models, NeRF, 3D Gaussian Splatting (3DGS)
Data & Systems:	Distributed Training, Multi-GPU Scaling, Large-scale Data Pipelines

Interests

- Reviewer for international conferences such as CVPR (2024, 2025), ECCV (2026) and MICCAI (2026).
- Passionate about basketball, cycling, and running.