

Stathis Galanakis

RESEARCH ENGINEER | MULTIMODAL GENERATIVE AI | MACHINE LEARNING

London, UK

@stathisgalan@gmail.com

in stathisgln |  Stathis Galanakis |  stathisgln.github.io |  stathisgln

Summary

Research Engineer (PhD, Imperial College London) specialising in generative AI, diffusion models, and multimodal machine learning systems. My work focuses on scalable generative modeling under low-data and real-world constraints, with applications spanning synthetic data generation, representation learning, and identity-consistent generation. Published first-author research at ICCV and WACV, with experience building large-scale datasets, training pipelines, and generative ML systems across academia and industry.

Experience

Research Associate in Generative Computer Vision Models

London, UK

Department of Computing, Imperial College London

March 2025 - Present

- Developed rectified flow-based image synthesis for low-data augmentation, resulting in an 8% improvement in pathology classification.
- Engineered LLM workflows for synthetic medical caption generation (500k image-text pairs) and explored automated pathology report classification.
- Supervised Master's students on research projects in generative AI and world modeling.

Research Engineer Intern

London, UK

Huawei UK

January 2022 - March 2025

- Researched diffusion-based generative models for robust identity-preserving image synthesis under ambiguous visual conditions.
- Curated synthetic identity datasets (10k+ identities) to improve model robustness and generalisation.
- Designed neural rendering pipelines for controllable multi-view avatar generation and consistent viewpoint synthesis.

Research Assistant

London, UK

Business School, Imperial College London

February 2021 - January 2022

- Applied machine learning to multimodal satellite imagery and weather data for crop yield prediction under spatial data sparsity.
- Developed synthetic data augmentation strategies to improve generalisation in underrepresented geographic regions.

Computer Vision Scientist

London, UK

ArielAI (acquired by Snap Inc.)

January 2020 - September 2020

- Built large-scale data collection and filtering pipelines for training vision models, incorporating human-in-the-loop annotation to improve dataset quality.
- Contributed to scalable data infrastructure supporting model development in a startup environment, later acquired by Snap Inc.

R&D, ML Engineer

Athens, Greece

Pobuca Ltd

May 2018 - January 2019

- Designed and deployed an end-to-end product detection system for retail shelf images (25+ classes), integrating semi-automated annotation and training workflows.

Selected Publications (7 total)

DERMAFLUX: Synthetic Skin Lesion Generation with Rectified Flows for Enhanced Image Classification
Galanakis S, Koliouris A, and Zafeiriou S. *Under peer review*
DERMAFLUX is a rectified flow-based text-to-image model for synthetic skin lesion image synthesis.

STARCaster: Spatio-Temporal AutoRegressive Video Diffusion for Identity- and View-Aware Talking Portraits
Paraperas F, Galanakis S, Potamias R, Kainz B, and Zafeiriou S. [ICML 2026](#)

SpinMeRound: Consistent Multi-View Identity Generation Using Diffusion Models
Galanakis S, Lattas A, Moschoglou S, Kainz B, and Zafeiriou S. [ICCV 2025](#)
SpinMeRound introduces a diffusion framework for identity-consistent multi-view visual generation from sparse inputs.

ImHead: A Large-scale Implicit Morphable Model for Localized Head Modeling
Potamias R, Galanakis S, Deng J, Papaioannou A, and Zafeiriou S. [ICCV 2025](#)

FitDiff: Robust monocular 3D facial shape and reflectance estimation using Diffusion Models
Galanakis S, Lattas A, Moschoglou S, and Zafeiriou S. [WACV 2025](#)
FitDiff introduces a multimodal diffusion framework for robust identity-preserving visual generation from unconstrained images.

3DMM-RF: Convolutional Radiance Fields for 3D Face Modeling
Galanakis S, Gecer B, Lattas A, and Zafeiriou S. [WACV 2023](#)
3DMM-RF explores neural radiance field representations for controllable and identity-consistent visual generation, under diverse expressions, poses and lighting conditions.

Education

PhD in Computer Science [Imperial College London, UK](#)
Supervisor: *Prof. Stefanos Zafeiriou* April 2025
Thesis: Advancements in Face Reconstruction from a Single Image.
Research focused on diffusion-based generative modeling and robust visual generation under unconstrained real-world conditions. Work resulted in peer-reviewed publications (ICCV, WACV), including first-author contributions.

Diploma/M.Eng. in Electrical and Computer Engineering [NTUA, Greece](#)
Supervisor: *Prof. Petros Maragos* November 2019
Thesis: Human Action Recognition and Localisation in Videos.

Technical Strengths

Programming:	Python, CUDA, C++
Frameworks:	PyTorch, PyTorch Lightning, PyTorch3D, Accelerate
Research Areas:	Diffusion Models, Rectified Flows, Transformers, GANs, 3D Gaussian Splatting (3DGS)
Data & Systems:	Distributed Training, Multi-GPU Scaling, Large-scale Data Pipelines

Interests

- Reviewer for international conferences such as CVPR (2024, 2025), ECCV (2026) and MICCAI (2026).
- Passionate about basketball, cycling, and running.