

Stathis Galanakis

RESEARCH ENGINEER | MULTIMODAL GENERATIVE AI | MACHINE LEARNING

London, UK

@stathisgalan@gmail.com | [in stathisgln](#) | [stathisgln.github.io](#) | [stathisgln](#) | [Stathis Galanakis](#)

Summary

Research Engineer (PhD, Imperial College London) specialising in generative AI, diffusion models, and multimodal machine learning systems. Published first-author work at ICCV/WACV, with experience building scalable multi-GPU training systems and synthetic data pipelines for low-data regimes.

Technical Strengths

Generative AI:	Diffusion Models, Video Diffusion, Rectified Flows, Transformers
3D & Vision:	3D Face Reconstruction, 3D Scene Reconstruction, NeRF, 3D Gaussian Splatting (3DGS)
Programming:	Python, CUDA, C++, Linux, Git, Docker
Frameworks:	PyTorch, PyTorch Lightning, PyTorch3D, Accelerate
Systems:	Distributed Training, Multi-GPU Scaling, SLURM, Large-scale Data Pipelines

Experience

Research Associate in Generative Computer Vision Models

London, UK

Department of Computing, Imperial College London

March 2025 - Present

- Developed rectified flow-based image synthesis for low-data augmentation, resulting in an 8% improvement in pathology classification.
- Scaled training of high-resolution generative models by deploying multi-GPU workflows (PyTorch DDP).
- Engineered LLM workflows for synthetic medical caption generation (500k image-text pairs) and explored automated pathology report classification.
- Supervised Master's students on research projects in generative AI and world modeling.

Research Engineer (Internship)

London, UK

Huawei UK

January 2022 - March 2025

- Researched diffusion-based generative models for robust identity-preserving image synthesis under ambiguous visual conditions.
- Curated synthetic identity datasets (10k+ identities) to improve model robustness and generalisation.
- Designed neural rendering pipelines for controllable multi-view avatar generation and consistent viewpoint synthesis.

Research Assistant

London, UK

Business School, Imperial College London

February 2021 - January 2022

- Developed machine learning models on 10k satellite imagery and weather data for agricultural crop yield prediction.
- Increased model generalisation in underrepresented areas by implementing synthetic data generation frameworks.

Computer Vision Scientist

London, UK

ArielAI (acquired by Snap Inc.)

January 2020 - September 2020

- Built large-scale data pipelines and human-in-the-loop annotation workflows, supporting model training at a startup acquired by Snap Inc.

R&D, ML Engineer

Athens, Greece

Pobuca Ltd

May 2018 - January 2019

- Delivered an end-to-end product detection system (25+ classes) for supermarket shelf images, improving in-store monitoring via semi-automated annotation workflows, training, and deployment.

Selected Publications (7 total)

SpinMeRound: Consistent Multi-View Identity Generation Using Diffusion Models

First Author - ICCV 2025

FitDiff: Robust monocular 3D facial shape and reflectance estimation using Diffusion Models

First Author - WACV 2025

STARCaster: Spatio-Temporal AutoRegressive Video Diffusion for Identity- and View-Aware Talking Portraits Co-Author - ICML 2026

DERMAFLUX: Synthetic Skin Lesion Generation with Rectified Flows for Enhanced Image Classification

First Author - MICCAI 2026

Education

PhD in Computer Science

Imperial College London, 2025

Diploma/M.Eng. in Electrical and Computer Engineering

NTUA, 2019