# Stathis **Galanakis**


Research Engineer | Multimodal Generative AI | Machine Learning

London, UK

@ stathisgalan@gmail.com | stathisgln | stathisgln.github.io | stathisgln | Stathis Galanakis


## Summary


Research Engineer (PhD, Imperial College London) specialising in generative AI, diffusion models, and multimodal machine learning systems. Published first-author work at ICCV/WACV, with experience building scalable multi-GPU training systems and synthetic data pipelines for low-data regimes.


## Technical Strengths

| | |
|---|---|
| **Generative AI:** | Diffusion Models, Video Diffusion, Rectified Flows, Transformers |
| **3D & Vision:** | 3D Face Reconstruction, 3D Scene Reconstruction, NeRF, 3D Gaussian Splatting (3DGS) |
| **Programming:** | Python, CUDA, C++, Linux, Git, Docker |
| **Frameworks:** | PyTorch, PyTorch Lightning, PyTorch3D, Accelerate |
| **Systems:** | Distributed Training, Multi-GPU Scaling, SLURM, Large-scale Data Pipelines |

## Experience

**Research Associate in Generative Computer Vision Models** — London, UK
*Department of Computing, Imperial College London* — *March 2025 - Present*

- Developed rectified flow-based image synthesis for low-data augmentation, resulting in an 8% improvement in pathology classification.
- Scaled training of high-resolution generative models by deploying multi-GPU workflows (PyTorch DDP).
- Engineered LLM workflows for synthetic medical caption generation (500k image-text pairs) and explored automated pathology report classification.
- Supervised Master's students on research projects in generative AI and world modeling.

**Research Engineer (Internship)** — London, UK
*Huawei UK* — *January 2022 - March 2025*

- Researched diffusion-based generative models for robust identity-preserving image synthesis under ambiguous visual conditions.
- Curated synthetic identity datasets (10k+ identities) to improve model robustness and generalisation.
- Designed neural rendering pipelines for controllable multi-view avatar generation and consistent viewpoint synthesis.

**Research Assistant** — London, UK
*Business School, Imperial College London* — *February 2021 - January 2022*

- Developed machine learning models on 10k satellite imagery and weather data for agricultural crop yield prediction.
- Increased model generalisation in underrepresented areas by implementing synthetic data generation frameworks.

**Computer Vision Scientist** — London, UK
*ArielAI (acquired by Snap Inc.)* — *January 2020 - September 2020*

- Built large-scale data pipelines and human-in-the-loop annotation workflows, supporting model training at a startup acquired by Snap Inc.

**R&D, ML Engineer** — Athens, Greece
*Pobuca Ltd* — *May 2018 - January 2019*

- Delivered an end-to-end product detection system (25+ classes) for supermarket shelf images, improving in-store monitoring via semi-automated annotation workflows, training, and deployment.

## Selected Publications (7 total)

**SpinMeRound: Consistent Multi-View Identity Generation Using Diffusion Models** — *First Author - ICCV 2025*

**FitDiff: Robust monocular 3D facial shape and reflectance estimation using Diffusion Models** — *First Author - WACV 2025*

**STARCaster: Spatio-Temporal AutoRegressive Video Diffusion for Identity- and View-Aware Talking Portraits** — ICML 2026

## Education

**PhD in Computer Science** — Imperial College London, 2025

**Diploma/M.Eng. in Electrical and Computer Engineering** — NTUA, 2019